

# Trust Region Policy Optimization Demo

Rishikesh Vaishnav

July 11, 2018

## Code

- The code for this project is available at: .

## Implementation Details

Single Path Executable Pseudocode:

- Initialize policy parameter  $\theta$ .
- Iterate until convergence:
  - Initialize/clear list  $S$  of  $\{\frac{G_\theta(s,a)}{\pi_\theta(s,a)}, s, a\}$ .
  - Generate  $N_\tau$  trajectories  $\{\tau\}$ .
  - For each trajectory  $\tau \in \{\tau\}$ :
    - For each  $\{s, a\} \in \tau$ :
      - Calculate discounted return  $G_\theta(s, a)$  from this time to end of episode.
      - Calculate  $\pi_\theta(s, a)$  at this  $(s, a)$ .
      - Store  $\{\frac{G_\theta(s,a)}{\pi_\theta(s,a)}, s, a\}$  in  $S$ .
  - Use constraint optimizer to yield  $\theta'$  by solving the problem:
    - Objective (to maximize):  $\text{objective}(S, \theta)$ .
    - Constraint:  $\text{constraint}(S, \theta)$ .
  - $\theta = \theta'$ .

$\text{objective}(S, \theta')$  Pseudocode:

- Initialize  $L = 0$ .
- For  $\{\frac{G_\theta(s,a)}{\pi_\theta(s,a)}, s, a\} \in S$ :
  - Add  $\pi_{\theta'}(s, a) \frac{G_\theta(s,a)}{\pi_\theta(s,a)}$  to  $L$ .
- Return  $L$ .

$\text{constraint}(S, \theta')$  Pseudocode:

- Initialize  $D = 0$ .
- For  $s \in S$ :

- Add  $D_{KL}(\pi_{\theta}(\cdot|s)||\pi_{\theta'}(\cdot|s))$  to  $D$ .
- Return  $\frac{D}{|S|}$ .

Parameter Settings:

- $N_{\tau} = 5$  (adjusted to balance between empirical runtime and performance)
- $\gamma = 1$  (adjusted to maximize empirical performance)
- $\delta = 0.01$  (following Schulman et. al.)

Policy Function Encoding:


- Each state-action pair is converted to the feature vector  $x(s, a)$ . Letting  $S_{obs}$  and  $S_{act}$  be the size of the observation and action spaces, respectively, the size of the vector is  $S_{obs} \times S_{act}$ , where all features are 0 except for the  $S_{obs}$  features starting at index  $S_{obs} \times a$ , which are set to the environment’s parameterization of  $s$ .
  - In this case,  $S_{obs} = 4$  and  $S_{act} = 2$ .
- The policy function  $\pi(a|s, \theta)$  performs the softmax on a parameterized linear mapping of feature vectors:

$$\pi(a|s, \theta) = \frac{e^{\theta^T x(s, a)}}{\sum_b e^{\theta^T x(s, b)}}$$

Constraint Optimization Method:

- I used scipy’s optimize.minimize function, with the “trust-constr” method. All gradients and hessians were automatically calculated. A tolerance of  $1 \times 10^{-2}$  and a maximum iteration limit of 500 were used.

## Results



- The results can be summarized as follows:

–