

Research Paper Ideas

Rishikesh Vaishnav

July 12, 2018

Trust Region Policy Optimization

Reducing Approximations

- How does the replacement of the penalty term with a constraint affect the monotonic improvement guarantee, and if the guarantee no longer holds, how can we (perhaps dynamically) adjust δ in the constraint to better ensure monotonic improvement?
 - This will likely require knowledge of constraint optimization algorithms.
- Instead of replacing D_{KL}^{max} with the expected KL-divergence \bar{D}_{KL}^ρ , can we replace it with something that better approximates D_{KL}^{max} ?
- Considering the theory behind equation (14) and the single path and vine implementations of TRPO, it seems like the authors implicitly approximate the normalized discounted visitation frequencies $(1 - \gamma)p_{\pi_{\theta_{old}}}(s) = (1 - \gamma)(P(s_0 = s) + \gamma P(s_1 = s) + \gamma^2 P(s_2 = s) + \dots)$ with the on-policy distribution of states under $\pi_{\theta_{old}}$. To what extent can this be justified? It seems that they are not necessarily equal, because the former depends on the discount factor γ , while the latter does not.

Improving Sample Efficiency

- To what extent can samples from previous runs (e.g. from runs using values of θ older than θ_{old}) be re-used in the current iteration? Perhaps importance sampling can help here?

Generalizations

- How can we generalize the vine method beyond simulated environments? Can resetting an uncontrolled environment to a state that is “similar enough” to the branch point ever yield competent performance?

Proximal Policy Optimization

Improving Clipped Approximation Function

- Is there a way to robustly choose a clipping parameter ϵ ? Can this parameter be dynamic?
- The sample advantage \hat{A}_t expanded in equation (10) seems inherently biased towards 0 because it makes use of a T -step return that bootstraps off of $V(s_T)$. How can this bias be addressed?

Investigating Other Potential Approximation Functions

- The clipping version of PPO only imposes a penalty (reduction in approximated policy value) for a given sign of \hat{A} if the policies diverge in one direction (where “direction” refers to whether $r_t(\theta) > 1$ or $r_t(\theta) < 1$), while the theory behind TRPO would suggest imposing a penalty in both directions. Only penalizing one direction could be helpful for increasing step size, but doesn’t this come at the cost of the monotonic improvement guarantee? Is there a more precise way of modulating the tradeoff between step size and monotonic improvement?
- The original paper on TRPO replaced the penalty term with a constraint because the penalty coefficient $C = \frac{4\epsilon\gamma}{(1-\gamma)^2}$ caused the step sizes to be “too small.” By what standard is “too small” defined, and, after determining what this is, can we alter β in equation (8) to make sure we are taking a step size that is as large as possible without being too big? Section 4 presents a technique for modifying β , but there doesn’t seem to be much theory behind it, so can that be improved?